

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [3]: df = pd.read_csv('7431_sales_data_sample.csv', encoding='unicode_escape')
```

```
In [5]: df.head()
```

```
Out[5]:    ORDERNUMBER QUANTITYORDERED PRICEEACH ORDERLINENUMBER SALES OR  
0          10107           30     95.70             2  2871.00  
1          10121           34     81.35             5  2765.90  
2          10134           41     94.74             2  3884.34  
3          10145           45     83.26             6  3746.70  
4          10159           49    100.00            14  5205.27
```

5 rows × 25 columns

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ORDERNUMBER      2823 non-null    int64  
 1   QUANTITYORDERED 2823 non-null    int64  
 2   PRICEEACH        2823 non-null    float64 
 3   ORDERLINENUMBER 2823 non-null    int64  
 4   SALES            2823 non-null    float64 
 5   ORDERDATE        2823 non-null    object  
 6   STATUS            2823 non-null    object  
 7   QTR_ID           2823 non-null    int64  
 8   MONTH_ID         2823 non-null    int64  
 9   YEAR_ID          2823 non-null    int64  
 10  PRODUCTLINE      2823 non-null    object  
 11  MSRP             2823 non-null    int64  
 12  PRODUCTCODE      2823 non-null    object  
 13  CUSTOMERNAME     2823 non-null    object  
 14  PHONE             2823 non-null    object  
 15  ADDRESSLINE1     2823 non-null    object  
 16  ADDRESSLINE2     302 non-null     object  
 17  CITY              2823 non-null    object  
 18  STATE             1337 non-null    object  
 19  POSTALCODE        2747 non-null    object  
 20  COUNTRY           2823 non-null    object  
 21  TERRITORY         1749 non-null    object  
 22  CONTACTLASTNAME   2823 non-null    object  
 23  CONTACTFIRSTNAME  2823 non-null    object  
 24  DEALSIZE          2823 non-null    object  
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

```
In [8]: #Columns to Remove
to_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATE', 'POSTALCODE', 'PHONE']
df = df.drop(to_drop, axis=1)
```

```
In [9]: #Check for null values
df.isnull().sum()
```

```
Out[9]: ORDERNUMBER          0
QUANTITYORDERED         0
PRICEEACH                0
ORDERLINENUMBER          0
SALES                     0
ORDERDATE                 0
STATUS                     0
QTR_ID                    0
MONTH_ID                  0
YEAR_ID                   0
PRODUCTLINE                0
MSRP                      0
PRODUCTCODE                0
CUSTOMERNAME               0
CITY                       0
COUNTRY                    0
TERRITORY                  1074
CONTACTLASTNAME            0
CONTACTFIRSTNAME            0
DEALSIZE                   0
dtype: int64
```

```
In [10]: df.dtypes
```

```
Out[10]: ORDERNUMBER        int64
QUANTITYORDERED        int64
PRICEEACH              float64
ORDERLINENUMBER        int64
SALES                  float64
ORDERDATE                object
STATUS                  object
QTR_ID                  int64
MONTH_ID                  int64
YEAR_ID                  int64
PRODUCTLINE                object
MSRP                     int64
PRODUCTCODE                object
CUSTOMERNAME               object
CITY                      object
COUNTRY                  object
TERRITORY                  object
CONTACTLASTNAME            object
CONTACTFIRSTNAME            object
DEALSIZE                  object
dtype: object
```

```
In [12]: #ORDERDATE Should be in date time
df['ORDERDATE'] = pd.to_datetime(df['ORDERDATE'])
```

```
In [13]: #We need to create some features in order to create cluseters
#Recency: Number of days between customer's latest order and today's date
#Frequency : Number of purchases by the customers
#MonetaryValue : Revenue generated by the customers
import datetime as dt
snapshot_date = df['ORDERDATE'].max() + dt.timedelta(days = 1)
df_RFМ = df.groupby(['CUSTOMERNAME']).agg({
    'ORDERDATE' : lambda x : (snapshot_date - x.max()).days,
    'ORDERNUMBER' : 'count',
    'SALES' : 'sum'}
```

```

    })

#Rename the columns
df_RFMs.rename(columns = {
    'ORDERDATE' : 'Recency',
    'ORDERNUMBER' : 'Frequency',
    'SALES' : 'MonetaryValue'
}, inplace=True)

```

In [14]: df\_RFMs.head()

Out[14]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue
<b>AV Stores, Co.</b>	196	51	157807.81
<b>Alpha Cognac</b>	65	20	70488.44
<b>Amica Models &amp; Co.</b>	265	26	94117.26
<b>Anna's Decorations, Ltd</b>	84	46	153996.13
<b>Atelier graphique</b>	188	7	24179.96

In [15]:

```

# Divide into segments
# We create 4 quartile ranges
df_RFMs['M'] = pd.qcut(df_RFMs['MonetaryValue'], q = 4, labels = range(1,5))
df_RFMs['R'] = pd.qcut(df_RFMs['Recency'], q = 4, labels = list(range(4,0,-1)))
df_RFMs['F'] = pd.qcut(df_RFMs['Frequency'], q = 4, labels = range(1,5))

df_RFMs.head()

```

Out[15]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue	M	R	F
<b>AV Stores, Co.</b>	196	51	157807.81	4	2	4
<b>Alpha Cognac</b>	65	20	70488.44	2	4	2
<b>Amica Models &amp; Co.</b>	265	26	94117.26	3	1	2
<b>Anna's Decorations, Ltd</b>	84	46	153996.13	4	3	4
<b>Atelier graphique</b>	188	7	24179.96	1	2	1

In [16]:

```

#Create another column for RFM score
df_RFMs['RFM_Score'] = df_RFMs[['R', 'M', 'F']].sum(axis=1)
df_RFMs.head()

```

Out[16]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue	M	R	F	RFM_Score
AV Stores, Co.	196	51	157807.81	4	2	4	10
Alpha Cognac	65	20	70488.44	2	4	2	8
Amica Models & Co.	265	26	94117.26	3	1	2	6
Anna's Decorations, Ltd	84	46	153996.13	4	3	4	11
Atelier graphique	188	7	24179.96	1	2	1	4

In [17]:

```
def rfm_level(df):
    if bool(df['RFM_Score'] >= 10):
        return 'High Value Customer'

    elif bool(df['RFM_Score'] < 10) and bool(df['RFM_Score'] >= 6):
        return 'Mid Value Customer'
    else:
        return 'Low Value Customer'
df_RFМ['RFM_Level'] = df_RFМ.apply(rfm_level, axis = 1)
df_RFМ.head()
```

Out[17]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue	M	R	F	RFM_Score	RFM_Level
AV Stores, Co.	196	51	157807.81	4	2	4	10	High Value Customer
Alpha Cognac	65	20	70488.44	2	4	2	8	Mid Value Customer
Amica Models & Co.	265	26	94117.26	3	1	2	6	Mid Value Customer
Anna's Decorations, Ltd	84	46	153996.13	4	3	4	11	High Value Customer
Atelier graphique	188	7	24179.96	1	2	1	4	Low Value Customer

In [18]:

```
# Time to perform KMeans
data = df_RFМ[['Recency', 'Frequency', 'MonetaryValue']]
data.head()
```

Out[18]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue
<b>AV Stores, Co.</b>	196	51	157807.81
<b>Alpha Cognac</b>	65	20	70488.44
<b>Amica Models &amp; Co.</b>	265	26	94117.26
<b>Anna's Decorations, Ltd</b>	84	46	153996.13
<b>Atelier graphique</b>	188	7	24179.96

In [19]:

```
# Our data is skewed we must remove it by performing Log transformation
data_log = np.log(data)
data_log.head()
```

Out[19]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue
<b>AV Stores, Co.</b>	5.278115	3.931826	11.969133
<b>Alpha Cognac</b>	4.174387	2.995732	11.163204
<b>Amica Models &amp; Co.</b>	5.579730	3.258097	11.452297
<b>Anna's Decorations, Ltd</b>	4.430817	3.828641	11.944683
<b>Atelier graphique</b>	5.236442	1.945910	10.093279

In [20]:

```
#Standardization
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(data_log)
data_normalized = scaler.transform(data_log)
data_normalized = pd.DataFrame(data_normalized, index = data_log.index, columns=
data_normalized.describe().round(2))
```

Out[20]:

	Recency	Frequency	MonetaryValue
<b>count</b>	92.00	92.00	92.00
<b>mean</b>	0.00	-0.00	0.00
<b>std</b>	1.01	1.01	1.01
<b>min</b>	-3.51	-3.67	-3.82
<b>25%</b>	-0.24	-0.41	-0.39
<b>50%</b>	0.37	0.06	-0.04
<b>75%</b>	0.53	0.45	0.52
<b>max</b>	1.12	4.03	3.92

In [21]:

```
#Fit KMeans and use elbow method to choose the number of clusters
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from sklearn.cluster import KMeans

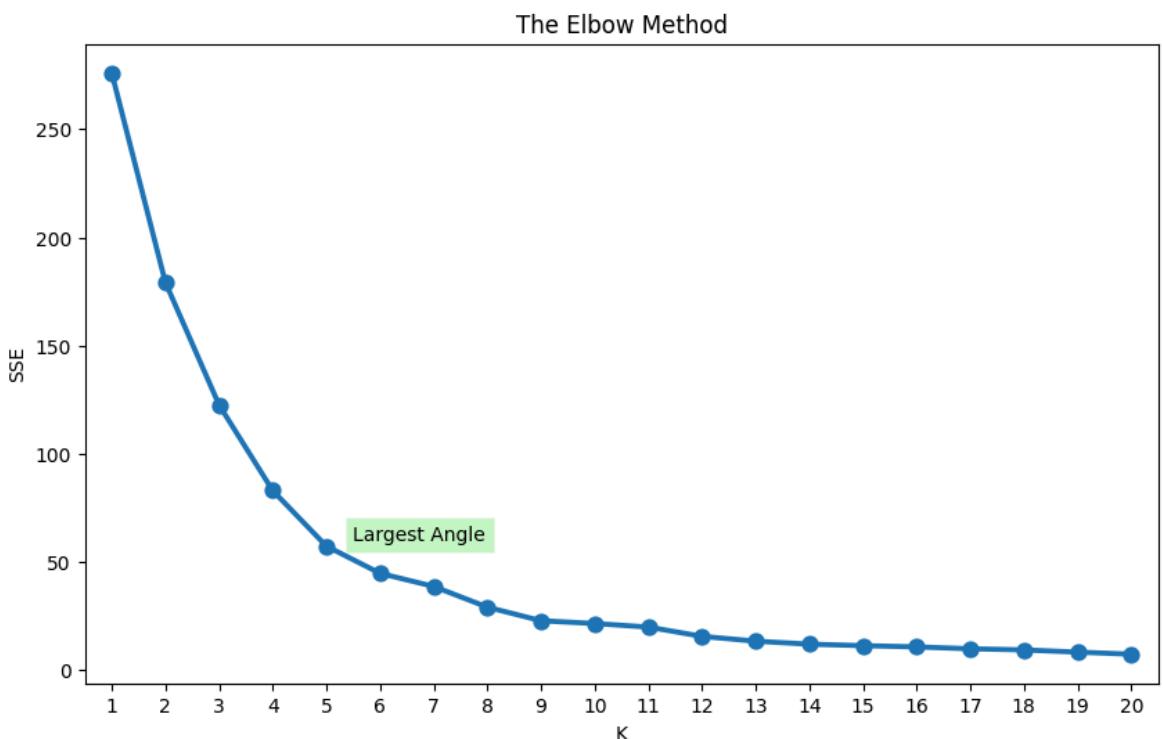
sse = {}

for k in range(1, 21):
    kmeans = KMeans(n_clusters = k, random_state = 1)
    kmeans.fit(data_normalized)
    sse[k] = kmeans.inertia_
```

```
In [22]: plt.figure(figsize=(10,6))
plt.title('The Elbow Method')

plt.xlabel('K')
plt.ylabel('SSE')
plt.style.use('ggplot')

sns.pointplot(x=list(sse.keys()), y = list(sse.values()))
plt.text(4.5, 60, "Largest Angle", bbox = dict(facecolor = 'lightgreen', alpha =
plt.show()
```



```
In [23]: # 5 number of clusters seems good
kmeans = KMeans(n_clusters=5, random_state=1)
kmeans.fit(data_normalized)
cluster_labels = kmeans.labels_

data_rf = data.assign(Cluster = cluster_labels)
data_rf.head()
```

Out[23]:

CUSTOMERNAME	Recency	Frequency	MonetaryValue	Cluster
<b>AV Stores, Co.</b>	196	51	157807.81	4
<b>Alpha Cognac</b>	65	20	70488.44	2
<b>Amica Models &amp; Co.</b>	265	26	94117.26	2
<b>Anna's Decorations, Ltd</b>	84	46	153996.13	4
<b>Atelier graphique</b>	188	7	24179.96	1