

```
In [3]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
#We do not want to see warnings
warnings.filterwarnings("ignore")
```

```
In [4]: #import data
data = pd.read_csv("7431_uber.csv")
```

```
In [5]: #Create a data copy
df = data.copy()
```

```
In [6]: #Print data
df.head()
```

Out[6]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pick
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	



```
In [7]: #Get Info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Unnamed: 0        200000 non-null   int64  
 1   key              200000 non-null   object  
 2   fare_amount      200000 non-null   float64 
 3   pickup_datetime  200000 non-null   object  
 4   pickup_longitude 200000 non-null   float64 
 5   pickup_latitude  200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude 199999 non-null   float64 
 8   passenger_count  200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
In [8]: #pickup_datetime is not in required data format
```

```
df["pickup_datetime"] = pd.to_datetime(df["pickup_datetime"])
```

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Unnamed: 0        200000 non-null   int64  
 1   key              200000 non-null   object  
 2   fare_amount       200000 non-null   float64 
 3   pickup_datetime   200000 non-null   datetime64[ns, UTC] 
 4   pickup_longitude  200000 non-null   float64 
 5   pickup_latitude   200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count   200000 non-null   int64  
dtypes: datetime64[ns, UTC](1), float64(5), int64(2), object(1)
memory usage: 13.7+ MB
```

In [10]: `#Statistics of data
df.describe()`

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude
count	2.000000e+05	200000.000000	200000.000000	200000.000000	199999.000000
mean	2.771250e+07	11.359955	-72.527638	39.935885	-72.525295
std	1.601382e+07	9.901776	11.437787	7.720539	13.117400
min	1.000000e+00	-52.000000	-1340.648410	-74.015515	-3356.666300
25%	1.382535e+07	6.000000	-73.992065	40.734796	-73.991400
50%	2.774550e+07	8.500000	-73.981823	40.752592	-73.980090
75%	4.155530e+07	12.500000	-73.967154	40.767158	-73.963650
max	5.542357e+07	499.000000	57.418457	1644.421482	1153.572600

In [11]: `#Number of missing values
df.isnull().sum()`

Unnamed: 0	0
key	0
fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	1
dropoff_latitude	1
passenger_count	0
dtype:	int64

In [12]: `#Correlation - only numeric columns
df.select_dtypes(include=[np.number]).corr()`

Out[12]:


	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_
Unnamed: 0	1.000000	0.000589	0.000230	-0.000341	
fare_amount	0.000589	1.000000	0.010457	-0.008481	
pickup_longitude	0.000230	0.010457	1.000000	-0.816461	
pickup_latitude	-0.000341	-0.008481	-0.816461	1.000000	
dropoff_longitude	0.000270	0.008986	0.833026	-0.774787	
dropoff_latitude	0.000271	-0.011014	-0.846324	0.702367	
passenger_count	0.002257	0.010150	-0.000414	-0.001560	



In [13]: `#Drop the rows with missing values
df.dropna(inplace=True)`

In [14]: `plt.boxplot(df['fare_amount'])`

Out[14]: {
 'whiskers': [`<matplotlib.lines.Line2D at 0x255a28602d0>`,
 `<matplotlib.lines.Line2D at 0x255a2862ad0>`],
 'caps': [`<matplotlib.lines.Line2D at 0x255a2862c10>`,
 `<matplotlib.lines.Line2D at 0x255a2862d50>`],
 'boxes': [`<matplotlib.lines.Line2D at 0x2559d596850>`],
 'medians': [`<matplotlib.lines.Line2D at 0x255a2862e90>`],
 'fliers': [`<matplotlib.lines.Line2D at 0x255a2862fd0>`],
 'means': []}



In [15]: `#Remove Outliers
q_low = df["fare_amount"].quantile(0.01)
q_hi = df["fare_amount"].quantile(0.99)`

```
df = df[(df["fare_amount"] < q_hi) & (df["fare_amount"] > q_low)]
```

In [16]: *#Check the missing values now*
`df.isnull().sum()`

Out[16]:

Unnamed:	0
key	0
fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	0
dropoff_latitude	0
passenger_count	0
dtype:	int64

In [17]: *#Time to apply Learning models*
`from sklearn.model_selection import train_test_split`

In [18]: *#Take x as predictor variable*
`x = df.drop("fare_amount", axis = 1)`
#And y as target variable
`y = df['fare_amount']`

In [19]: *#Necessary to apply model*
`x['pickup_datetime'] = pd.to_numeric(pd.to_datetime(x['pickup_datetime']))`
`x = x.loc[:, x.columns.str.contains('^Unnamed')]`

In [20]: `x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 42)`

In [21]: `from sklearn.linear_model import LinearRegression`

In [22]: `lrmrmodel = LinearRegression()`
`lrmrmodel.fit(x_train, y_train)`

Out[22]:

▼ LinearRegression ⓘ ⓘ

▼ Parameters

fit_intercept	True
copy_X	True
tol	1e-06
n_jobs	None
positive	False

In [23]: *#Prediction*
`predict = lrmrmodel.predict(x_test)`

In [24]: *#Check Error*
`from sklearn.metrics import mean_squared_error`

```
lrmode1rmse = np.sqrt(mean_squared_error(predict, y_test))
print("RMSE error for the model is ", lrmode1rmse)
```

RMSE error for the model is 8.063863046328837

In [25]: #Let's Apply Random Forest Regressor

```
from sklearn.ensemble import RandomForestRegressor
rfrmmodel = RandomForestRegressor(n_estimators = 100, random_state = 101)
```

In [26]: #Fit the Forest

```
rfrmmodel.fit(x_train, y_train)
rfrmmodel_pred = rfrmmodel.predict(x_test)
```

In [27]: #Errors for the forest

```
rfrmmodel_rmse = np.sqrt(mean_squared_error(rfrmmodel_pred, y_test))
print("RMSE value for Random Forest is:",rfrmmodel_rmse)
```

RMSE value for Random Forest is: 9.757713738069647